

# Stochastic Gradient Descent

## STAT 672 Project

Tom Wallace

George Mason University


Spring 2018

# Optimization is everywhere, and sometimes is easy

Many statistical procedures involve minimizing or maximizing some function applied to data

In **parametric** statistics, we often make assumptions that make this optimization “nice”

- Example: in OLS, we do not need to try different values of  $\hat{\beta}$  to see which minimizes the loss function, we (typically) can just evaluate  $(X'X)^{-1}X'Y$



# Other times, optimization is not so easy


Suppose that we have a typical supervised classification problem

- Non-parametric: no assumptions about distribution of data
- Feature vector  $\mathbf{X}_i$ , label  $Y_i$
- Want to find best prediction function  $f_w^*$  from class  $\mathcal{F}$
- Optimization: pick weights  $\mathbf{w}$  that minimize empirical risk according to some convex loss function  $L(f_w(x), y)$

Our lack of assumptions requires a different approach to optimization

- Cannot analytically identify stationary point
- Need to numerically search for it

# Gradient descent is an iterative search procedure



# A more formal explanation

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \gamma \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{w}} L(f_{\mathbf{w}}(\mathbf{X}_i), y_i)$$

Stop if gradient  $\leq \epsilon$

Step size  $\gamma$  can vary over time

Theoretical guarantees on speed of convergence ( $\rho :=$  size of error):

- Version presented here:  $-\log \rho \sim t$
- More optimized version:  $-\log \log \rho \sim t$

But, big difference between:

- Speed  $:=$  number of iterations
- Speed  $:=$  time (clock on wall)

# Batch gradient descent is computationally expensive

In “plain” (batch) gradient descent, for **every step**, we have to evaluate the gradient at **every observation**

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \gamma \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{w}} L(f_{\mathbf{w}}(\mathbf{x}_i), y_i)$$

This becomes computationally intractable as  $n$  grows large

Knowing that the quality of our approximation gets linearly or quadratically better with  $t$  is not comforting if each  $t$  takes days to run

# Stochastic gradient descent (SGD) takes less time

For each step, gradient is computed for a **single** randomly chosen observation  $i$ :

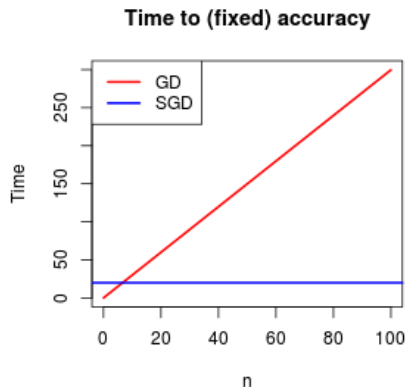
$$\mathbf{w}_{t+1} = \mathbf{w}_t - \gamma \nabla_w L(f_w(\mathbf{X}_i), y_i)$$

This simplification makes approximation much “noisier”, and hence SGD requires more iterations

But, each iteration is faster and so SGD can reach a predefined level of risk or error in less time

SGD is particularly useful when  $n$  is large and computation time is important

	GD	SGD
Time to accuracy $\rho$	$n \log \frac{1}{\rho}$	$\frac{1}{\rho}$





# SGD is widely used in industry

If a Silicon Valley press release uses any of the following phrases...

- “Neural networks”
- “Machine learning”
- “AI”

...SGD probably is involved. Example: Google’s **AlphaGo** program.

