Stochastic Gradient Descent

STAT 672 Project

Tom Wallace

George Mason University

Spring 2018

Estimating model coefficients requires optimization

We typically fit statistical models to maximize some measure of goodness (e.g., likelihood) or minimize some measure of badness (e.g., risk)

Parametric assumptions can make this optimization very nice

OLS has a closed form solution: $\hat{\beta} = (\mathbf{X}'\mathbf{X})\mathbf{X}'\mathbf{Y}$

GLM often uses Newton's method

- Knowing the PDF makes finding the Hessian (2nd order derivatives) not too onerous

Non-parametric & high-dimensional settings

Suppose that we have a typical supervised classification problem

- Non-parametric: no assumptions about distribution of data
- Feature vector $\mathbf{X}_i$, label Y_i
- Want to find best prediction function $f^*(\mathbf{w}; \mathbf{X})$ from hypothesis class $\mathcal{F}$
- Optimization: pick weights $\hat{\mathbf{w}}$ that minimize empirical risk according to some convex loss function

Some familiar tools for finding $\hat{\mathbf{w}}$ no longer work well

But we still know that if L is convex, there is a unique global minimum, and the gradient at that point must be 0

Gradient descent is a numerical strategy to search for that point

Gradient descent is an iterative search procedure

A more formal explanation

$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \frac{1}{n} \sum_{i=1}^{n} \nabla_w L(f(\mathbf{w}^{(t)}; \mathbf{X}_i), Y_i)$

Stop if gradient $\leq \epsilon$

Step size γ can vary over time

Batch gradient descent is computationally expensive

In "plain" (batch) gradient descent, for every step, we have to evaluate the gradient at every observation

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \frac{1}{n} \sum_{i=1}^{n} \nabla_w L(f(\mathbf{w}^{(t)}; \mathbf{X}_i), Y_i)$$

This becomes computationally intractable as n grows large

- If $n = 10$ million, have to evaluate gradient 10 million times every step

Stochastic gradient descent (SGD) takes less time

For each step, gradient is computed for a single randomly chosen observation i:

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \nabla_w L(f(\mathbf{w}^{(w)}; \mathbf{X}_i), Y_i)$$

This simplification makes approximation much "noisier", and hence SGD requires more iterations

But, each iteration is faster and so SGD can reach a predefined level of risk or error in less time

SGD is noisier than batch GD

Batch gradient descent
Mini-batch gradient Descent
Stochastic gradient descent

Source: towardsdatascience.com

SGD is useful when *n* is large & compute time is important

	GD	**SGD**
Time to accuracy ρ	$n \log \frac{1}{\rho}$	$\frac{1}{\rho}$

Figure: Fitting OLS coefficients, $n = 10$ million

SGD is very popular in academia and industry

Many, many statistical learning models use SGD to fit weights

If a Silicon Valley press release uses any of the following phrases...

- "Neural networks"
- "Machine learning"
- "AI"

...SGD probably is involved.

Example: Google's **AlphaGo** program.

Supervised learning of policy networks

Policy network: 12 layer convolutional neural network

Training data: 30M positions from human expert games (KGS 5+ dan)

Training algorithm: maximise likelihood by stochastic gradient descent

$$\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma}$$

Training time: 4 weeks on 50 GPUs using Google Cloud

Results: 57% accuracy on held out test data (state-of-the art was 44%)

Google DeepMind